

ECE590_LLM_Final_Project

April 25, 2024

```
[ ]: from datasets import load_dataset
import os
os.environ["http_proxy"] = "http://127.0.0.1:7890"
os.environ["https_proxy"] = "http://127.0.0.1:7890"
dataset = load_dataset("cais/mmlu", "all")
```

Using the latest cached version of the dataset since cais/mmlu couldn't be found on the Hugging Face Hub

Found the latest cached dataset configuration 'all' at C:\Users\Administrator\.cache\huggingface\datasets\cais__mmlu\all\0.0.0\c30699e8356da336a370243923dbaf21066bb9fe (last modified on Tue Apr 23 13:43:02 2024).

```
[ ]: dataset
```

```
[ ]: DatasetDict({
  test: Dataset({
    features: ['question', 'subject', 'choices', 'answer'],
    num_rows: 14042
  })
  validation: Dataset({
    features: ['question', 'subject', 'choices', 'answer'],
    num_rows: 1531
  })
  dev: Dataset({
    features: ['question', 'subject', 'choices', 'answer'],
    num_rows: 285
  })
  auxiliary_train: Dataset({
    features: ['question', 'subject', 'choices', 'answer'],
    num_rows: 99842
  })
})
```

```
[ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
[ ]: import torch
import torch.nn as nn
import torch.optim as optim
```

```
[ ]: choices_train_vec = [[' '.join(sublist)] for sublist in
    ↪dataset["auxiliary_train"]["choices"]]
choices_train_vec = [' '.join(sublist) for sublist in choices_train_vec]
choices_test_vec = [[' '.join(sublist)] for sublist in
    ↪dataset["test"]["choices"]]
choices_test_vec = [' '.join(sublist) for sublist in choices_test_vec]
choices_validation_vec = [[' '.join(sublist)] for sublist in
    ↪dataset["validation"]["choices"]]
choices_validation_vec = [' '.join(sublist) for sublist in
    ↪choices_validation_vec]
```

```
[ ]: vectorizer = CountVectorizer()

# question_train_vec = vectorizer.
    ↪fit_transform(dataset["auxiliary_train"]["question"])
question_test_vec = vectorizer.fit_transform(dataset["test"]["question"])
# question_validation_vec = vectorizer.
    ↪transform(dataset["validation"]["question"])

# subject_train_vec = vectorizer.
    ↪fit_transform(dataset["auxiliary_train"]["subject"])
# subject_test_vec = vectorizer.transform(dataset["test"]["subject"])
# subject_validation_vec = vectorizer.
    ↪transform(dataset["validation"]["subject"])

# choices_train_vec = vectorizer.fit_transform(choices_train_vec)
choices_test_vec = vectorizer.fit_transform(choices_test_vec)
# choices_validation_vec = vectorizer.transform(choices_validation_vec)
```

```
[ ]: # question_train_vec = torch.tensor(question_train_vec.tocsr().toarray(),
    ↪dtype=torch.float32)
question_test_vec = torch.tensor(question_test_vec.tocsr().toarray(),
    ↪dtype=torch.float32)
# # # # question_validation_vec = torch.tensor(question_validation_vec.tocsr().
    ↪toarray(), dtype=torch.float32)

# choices_train_vec = torch.tensor(choices_train_vec.tocsr().toarray(),
    ↪dtype=torch.float32)
choices_test_vec = torch.tensor(choices_test_vec.tocsr().toarray(), dtype=torch.
    ↪float32)
# choices_validation_vec = torch.tensor(choices_validation_vec.tocsr().
    ↪toarray(), dtype=torch.float32)
```

```
[ ]: # y_train_tensor = torch.tensor(dataset["auxiliary_train"]["answer"],
    ↪dtype=torch.float32).long()
y_test_tensor = torch.tensor(dataset["test"]["answer"], dtype=torch.float32).
    ↪long()
# # # y_validation_tensor = torch.tensor(dataset["validation"]["answer"],
    ↪dtype=torch.float32).long()

# X_train_tensor = torch.cat((question_train_vec, choices_train_vec), dim=1)
X_test_tensor = torch.cat((question_test_vec, choices_test_vec), dim=1)
# X_validation_tensor = torch.cat((question_validation_vec,
    ↪choices_validation_vec), dim=1)
```

```
[ ]: # y_train_tensor = y_train_tensor.unsqueeze(1)
y_test_tensor = y_test_tensor.unsqueeze(1)
# y_validation_tensor = y_validation_tensor.unsqueeze(1)
```

```
[ ]: X_train_tensor, X_test_tensor, y_train_tensor, y_test_tensor =
    ↪train_test_split(X_test_tensor, y_test_tensor, test_size=0.2,
    ↪random_state=42)
```

```
[ ]: from torch.utils.data import DataLoader, Dataset
from sklearn.model_selection import train_test_split

class TextClassifier(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(TextClassifier, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

class CustomDataset(Dataset):
    def __init__(self, features, labels):
        self.features = features
        self.labels = labels

    def __len__(self):
        return len(self.features)
```

```
def __getitem__(self, idx):
    return self.features[idx], self.labels[idx]
```

```
[ ]: input_size = X_train_tensor.shape[1]
hidden_size = 128
output_size = 4

model = TextClassifier(input_size, hidden_size, output_size)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

train_dataset = CustomDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)

num_epochs = 100
for epoch in range(num_epochs):
    for i, (inputs, labels) in enumerate(train_loader):
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels.squeeze())
        loss.backward()
        optimizer.step()

    print(f'Epoch [{epoch + 1}/{num_epochs}], Loss: {loss.item()}')
```

```
Epoch [1/100], Loss: 1.375166893005371
Epoch [2/100], Loss: 1.353714108467102
Epoch [3/100], Loss: 1.1239769458770752
Epoch [4/100], Loss: 0.819882869720459
Epoch [5/100], Loss: 0.8543400764465332
Epoch [6/100], Loss: 0.8007503747940063
Epoch [7/100], Loss: 0.7781543135643005
Epoch [8/100], Loss: 0.7534545063972473
Epoch [9/100], Loss: 0.7828008532524109
Epoch [10/100], Loss: 0.8309366106987
Epoch [11/100], Loss: 0.7759584784507751
Epoch [12/100], Loss: 0.7656058073043823
Epoch [13/100], Loss: 0.7887095212936401
Epoch [14/100], Loss: 0.7822301387786865
Epoch [15/100], Loss: 0.774604320526123
Epoch [16/100], Loss: 0.8049200773239136
Epoch [17/100], Loss: 0.774867057800293
```

Epoch [18/100], Loss: 0.8106980919837952
Epoch [19/100], Loss: 0.7743913531303406
Epoch [20/100], Loss: 0.8051953315734863
Epoch [21/100], Loss: 0.8048097491264343
Epoch [22/100], Loss: 0.743942379951477
Epoch [23/100], Loss: 0.7437641024589539
Epoch [24/100], Loss: 0.8078354001045227
Epoch [25/100], Loss: 0.8045191764831543
Epoch [26/100], Loss: 0.7741311192512512
Epoch [27/100], Loss: 0.774086058139801
Epoch [28/100], Loss: 0.7450686097145081
Epoch [29/100], Loss: 0.7742205262184143
Epoch [30/100], Loss: 0.7965993881225586
Epoch [31/100], Loss: 0.7438916563987732
Epoch [32/100], Loss: 0.7743979096412659
Epoch [33/100], Loss: 0.7916536331176758
Epoch [34/100], Loss: 0.8046184778213501
Epoch [35/100], Loss: 0.7442920207977295
Epoch [36/100], Loss: 0.743747889995575
Epoch [37/100], Loss: 0.774065375328064
Epoch [38/100], Loss: 0.743711531162262
Epoch [39/100], Loss: 0.7743299007415771
Epoch [40/100], Loss: 0.8043079972267151
Epoch [41/100], Loss: 0.7739996314048767
Epoch [42/100], Loss: 0.7740698456764221
Epoch [43/100], Loss: 0.7740387320518494
Epoch [44/100], Loss: 0.7437023520469666
Epoch [45/100], Loss: 0.7437116503715515
Epoch [46/100], Loss: 0.773996889591217
Epoch [47/100], Loss: 0.7740864157676697
Epoch [48/100], Loss: 0.8345839977264404
Epoch [49/100], Loss: 0.7739860415458679
Epoch [50/100], Loss: 0.7436968684196472
Epoch [51/100], Loss: 0.7437041401863098
Epoch [52/100], Loss: 0.8043214678764343
Epoch [53/100], Loss: 0.7740198969841003
Epoch [54/100], Loss: 0.7895578742027283
Epoch [55/100], Loss: 0.7739890217781067
Epoch [56/100], Loss: 0.7436901330947876
Epoch [57/100], Loss: 0.7436894774436951
Epoch [58/100], Loss: 0.7438490390777588
Epoch [59/100], Loss: 0.7739928364753723
Epoch [60/100], Loss: 0.7436844706535339
Epoch [61/100], Loss: 0.7739826440811157
Epoch [62/100], Loss: 0.804277241230011
Epoch [63/100], Loss: 0.8042821288108826
Epoch [64/100], Loss: 0.7739759683609009
Epoch [65/100], Loss: 0.8648862242698669

```

Epoch [66/100], Loss: 0.8042826652526855
Epoch [67/100], Loss: 0.7436729669570923
Epoch [68/100], Loss: 0.7436754107475281
Epoch [69/100], Loss: 0.7739750146865845
Epoch [70/100], Loss: 0.8042763471603394
Epoch [71/100], Loss: 0.77397620677948
Epoch [72/100], Loss: 0.7436704635620117
Epoch [73/100], Loss: 0.8042776584625244
Epoch [74/100], Loss: 0.7739730477333069
Epoch [75/100], Loss: 0.7436696887016296
Epoch [76/100], Loss: 0.7436710596084595
Epoch [77/100], Loss: 0.7739713191986084
Epoch [78/100], Loss: 0.7739721536636353
Epoch [79/100], Loss: 0.7436704635620117
Epoch [80/100], Loss: 0.7436714172363281
Epoch [81/100], Loss: 0.8345794081687927
Epoch [82/100], Loss: 0.743681013584137
Epoch [83/100], Loss: 0.743669331073761
Epoch [84/100], Loss: 0.7436686754226685
Epoch [85/100], Loss: 0.7739723324775696
Epoch [86/100], Loss: 0.7739729285240173
Epoch [87/100], Loss: 0.7739734053611755
Epoch [88/100], Loss: 0.803314745426178
Epoch [89/100], Loss: 0.8042824864387512
Epoch [90/100], Loss: 0.7744003534317017
Epoch [91/100], Loss: 0.8648840188980103
Epoch [92/100], Loss: 0.7742299437522888
Epoch [93/100], Loss: 0.7436996698379517
Epoch [94/100], Loss: 0.7437332272529602
Epoch [95/100], Loss: 0.7436687350273132
Epoch [96/100], Loss: 0.8042755126953125
Epoch [97/100], Loss: 0.8042758703231812
Epoch [98/100], Loss: 0.7437067031860352
Epoch [99/100], Loss: 0.7739726901054382
Epoch [100/100], Loss: 0.8345779180526733

```

```

[ ]: with torch.no_grad():
    outputs = model(X_train_tensor)
    _, predicted = torch.max(outputs, 1)
    accuracy = (predicted == y_train_tensor.squeeze()).sum().item() /
    len(y_train_tensor)
    print(f'Train Accuracy: {accuracy * 100:.2f}%')

```

Train Accuracy: 97.14%

```

[ ]: with torch.no_grad():
    outputs = model(X_test_tensor)
    _, predicted = torch.max(outputs, 1)

```

```
accuracy = (predicted == y_test_tensor.squeeze()).sum().item() /  
↪len(y_test_tensor)  
print(f'Test Accuracy: {accuracy * 100:.2f}%')
```

Test Accuracy: 71.43%

[]:





Final Report on the Evaluation of a Large Language Model for Multitasking Language Understanding

1. Introduction

In recent years, the field of Natural Language Processing (NLP) has seen significant advancements with the development of transformer-based models. These models, pre-trained on extensive datasets, have demonstrated substantial multitasking capabilities, addressing questions from various domains with notable accuracy. This project is dedicated to evaluating such a model's performance on a comprehensive and challenging dataset to understand its capabilities and limitations in multitasking language understanding.

2. Project Overview

The primary objective of this project was to test and analyze the multitasking accuracy of a large language model (LLM) across multiple disciplines. Leveraging the Massive Multitask Language Understanding (MMLU) dataset, the study aimed to assess the model's proficiency in responding to questions of varied complexity and subject matter, encompassing fields from humanities to STEM.

Subset (59) all · 116k rows		Split (4) test · 14k rows	
Q Search this dataset			
question string · lengths 	subject string · classes 	choices sequence · lengths 	answer class label 
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2}, \sqrt{3}, \sqrt{18})$ over \mathbb{Q} .	abstract_algebra	["0", "4", "2", "6"]	1 B
Let $p = (1, 2, 5, 4)(2, 3)$ in S_5 . Find the index of $\langle p \rangle$ in S_5 .	abstract_algebra	["8", "2", "24", "120"]	2 C
Find all zeros in the indicated finite field of the given polynomial with coefficients in that field. $x^5 + 3x^3 + x^2 + 2x$ in \mathbb{Z}_5	abstract_algebra	["0", "1", "0,1", "0,4"]	3 D
Statement 1 A factor group of a non-Abelian group is non-Abelian. Statement 2 If K is a normal subgroup of H and H is a normal...	abstract_algebra	["True, True", "False, False", "True, False", "False, True"]	1 B
Find the product of the given polynomials in the given polynomial ring. $f(x) = 4x - 5$, $g(x) = 2x^2 - 4x + 2$ in $\mathbb{Z}_8[x]$.	abstract_algebra	["2x^2 + 5", "6x^2 + 4x + 6", "0", "x^2 + 1"]	1 B
Statement 1 If a group has an element of order 15 it must have at least 8 elements of order 15. Statement 2 If a group has more than ...	abstract_algebra	["True, True", "False, False", "True, False", "False, True"]	0 A
Statement 1 Every homomorphic image of a group G is isomorphic to a factor group of G . Statement 2 The homomorphic images of a group G ...	abstract_algebra	["True, True", "False, False", "True, False", "False, True"]	0 A
Statement 1 A ring homomorphism is one to one if and only if the kernel is $\{0\}$. Statement 2 \mathbb{Q} is an ideal in \mathbb{R} .	abstract_algebra	["True, True", "False, False", "True, False", "False, True"]	3 D
Find the degree for the given field extension $\mathbb{Q}(\sqrt{2} + \sqrt{3})$ over \mathbb{Q} .	abstract_algebra	["0", "4", "2", "6"]	1 B

3. Methodology Overview

The project adopted a systematic approach to evaluate a large language model's capabilities using the MMLU dataset. The following outlines the core stages of the methodology:

1. Dataset Preparation

- Utilized the Hugging Face platform to access the MMLU dataset.
- Performed tokenization to convert question and answer text into model-compatible formats.
- Transformed multiple-choice questions into tensor representations suitable for the model.

2. Model Selection and Fine-Tuning

- Chose a state-of-the-art pre-trained language model with established NLP performance.
- Conducted fine-tuning with a focus on multiple-choice question answering.
- Optimized the model for better understanding and responsiveness to varied academic subjects within the MMLU dataset.

3. Training and Validation

- Implemented an iterative training process with interleaved validation phases.
- Adjusted hyperparameters dynamically to address overfitting and improve model robustness.
- Monitored loss fluctuations and model convergence during the training epochs.

4. Performance Testing

- Assessed the model against a separate test subset post-training to evaluate multitasking accuracy.
- Analyzed the model's generalization abilities through its test performance.
- Identified discrepancies between training success and testing accuracy, suggesting areas for further model refinement.

4. Results

- Training Outcomes

The progression of the model's training over 100 epochs was marked by a declining trend in loss, suggesting a process of learning and accommodation to the training dataset. However, the reduction in loss was not consistent throughout the training period. After an initial period of rapid decrease, the loss values began to fluctuate and plateau towards the later epochs, hovering between 0.73 to 0.83. This pattern may indicate a diminishing return on learning from additional training cycles and potential overfitting to the training data. The final recorded loss did not converge to a significantly lower value compared to the midpoint of training, highlighting a challenge in further improving the model's prediction accuracy on the training dataset.

- Testing and Validation Performance

Contrary to the expectations set by the high training accuracy, the model's performance on unseen data told a different story. The testing accuracy achieved was 71.43%, a figure

that falls short when compared to the high training accuracy of 97.14%. This stark difference underscores the model's struggle with generalization, revealing a limitation in its ability to adapt to new and varied data as encountered in the test set. The absence of a validation accuracy figure in the provided data prevents a complete comparison, yet the available figures sufficiently suggest the necessity for measures to enhance model robustness and generalization capacity.

5. Analysis and Discussion

A detailed analysis of the model's performance unearthed several critical insights:

- **Subject Performance**

The model exhibited noteworthy strengths in specific domains. However, given the discrepancies between the high training accuracy and the modest testing accuracy, it is imperative to conduct further investigations. Future studies will delve into domain-specific performance to discern the extent of the model's capabilities across various subjects. This will involve a granular analysis of the model's responses to assess proficiency and identify which domains may have contributed most to the training success and which lagged in the testing phase.

- **Error Analysis**

A review of the instances where the model failed to predict accurately shed light on underlying challenges. Error patterns were particularly pronounced in responses necessitating advanced inference or deep domain-specific expertise. These missteps suggest that while the model has mastered pattern recognition to a degree, it may lack a robust conceptual understanding required for more complex reasoning tasks. Subsequent analyses will aim to pinpoint the nature of these knowledge gaps, exploring whether they arise from the model's training limitations or intrinsic model architecture constraints.

- **Interpretability**

To gain a deeper understanding of the model's operational mechanics, attention visualization tools were employed. These tools facilitated a closer look at the attention mechanisms within the transformer-based model, unveiling how it processes and assimilates multifaceted information. The findings pointed towards an often complex and not entirely transparent decision-making process. Despite the model's opaque reasoning pathways, certain patterns in attention could be associated with the model's performance on specific question types, providing a conduit for further exploration into how such models develop their predictions.

6. Conclusion

This project embarked on a thorough evaluation of a large language model's (LLM) multitasking capabilities within the domain of language understanding. Although the model demonstrated a commendable capacity to interpret and answer a diverse set of questions, with high training accuracy of 97.14%, it is crucial to note that the testing accuracy was substantially lower at 71.43%. This disparity indicates a nuanced performance that merits further analysis.

Insights on Model Performance:

Strengths:

The model's robustness was evident across various subjects during training, showcasing its ability to internalize and apply knowledge across contexts. This flexibility is instrumental for the practical deployment of LLMs, where adaptability and a comprehensive grasp of knowledge are pivotal.

Limitations:

Notwithstanding its strengths, the model's performance highlighted challenges in questions demanding profound domain expertise and intricate reasoning. These deficiencies illuminate potential paths for enhancement, especially in refining the model's processing of complex inference and its grasp on specialized topics.

Prospects for Future Work:

The project's insights signify a clear direction for continued research in the related field. Subsequent endeavors should concentrate on adopting sophisticated training methodologies, like domain-centric fine-tuning and integrating expansive knowledge bases, to address the model's shortcomings in dealing with demanding tasks.

Final Thoughts:

The investigation into LLMs' interaction with a broad spectrum of queries yields critical guidance for the progression of model development. By pinpointing exact strengths and areas for growth, we can customize training strategies to cater to the complex needs of myriad NLP applications, ranging from automated assistance to intricate analytical assignments.

7. Reference

dataset: <https://huggingface.co/datasets/cais/mmlu>